



ALLIANCE
UNIVERSITY

*Private University established in Karnataka State by Act No.34 of year 2010
Recognized by the University Grants Commission (UGC), New Delhi*

ALLIANCE UNIVERSITY

Alliance School of Advanced Computing

Student Name: Shaurya Venkta

Register Number: PROV/ASAC/MCA/26/07/015

Project Name: BiteWise(Cloud Kitchen Menu Redundancy & Trade-Area Whitespace Prediction System)

ACADEMIC YEAR: 2026-2028

Department of Computer Applications

Program - Master of Computer Applications

Semester 1



ACKNOWLEDGEMENT

I would like to express my sincere gratitude to everyone who supported and guided me throughout the successful completion of this project.

I am thankful to my Project Guide for providing valuable guidance, suggestions, encouragement, and continuous support throughout the development of this project. I also express my gratitude to the faculty members and my institution for providing the necessary resources and a supportive learning environment.

I would also like to thank my friends and classmates for their cooperation and support during the completion of this project. Finally, I am deeply grateful to my family for their constant encouragement, motivation, and support.

I sincerely thank everyone who contributed directly or indirectly to the successful completion of this project.

Student Name: Shaurya Venkta

Register Number: PROV/ASAC/MCA/26/07/015

Project Title: BiteWise(Cloud Kitchen Menu Redundancy & Trade-Area Whitespace Prediction System)

Signature of Student

Signature of Faculty & Project Guide

Project Faculty & Project Guide Name: Oleti Durga Bravish

1. Abstract:

The exponential proliferation of on-demand online food delivery aggregators—most prominently Zomato and Swiggy—has fundamentally restructured the urban culinary economy, catalyzing the rapid emergence of cloud kitchens. Alternatively designated as ghost, virtual, or dark kitchens, these commercial food preparation facilities operate entirely without physical dine-in storefronts, relying exclusively on digital catalog visibility, third-party logistics fleets, and mobile ordering applications. While this operating paradigm dramatically mitigates traditional capital expenditures related to prime urban real estate, interior architectural design, and front-of-house hospitality personnel, it introduces acute, unquantified digital shelf competition within concentrated geographic delivery perimeters. Across major metropolitan food consumption corridors in India, particularly within dense delivery clusters across Bangalore and Delhi (such as Koramangala, Indiranagar, Saket, and Connaught Place), culinary operators encounter alarming commercial failure rates exceeding sixty percent within their initial six months of operation. A primary causal factor driving these business liquidations is extreme category cannibalization and unchecked menu redundancy: aspiring operators routinely launch duplicate culinary offerings—frequently characterized by undifferentiated biryani, burger, pizza, and generic North Indian curry concepts—into trade localities that are already hyper-saturated with statistically indistinguishable offerings, while lucrative, unmet consumer demand gaps and dietary niches remain entirely unaddressed.

To resolve this critical empirical challenge, this project introduces the design, mathematical formulation, and end-to-end implementation of an Applied Data Science (ADS) and Machine Learning system that systematically quantifies micro-trade-area menu redundancy and pinpoints commercially viable market whitespace. The foundation of this system rests on a multi-source data fusion methodology that integrates five heterogeneous empirical datasets: a web-scraped delivery catalog comprising over 4,000 dish-level records from Bangalore and Delhi with price bands and ratings; a curated culinary taxonomy derived from the Indian Food 101 database encompassing 255 ground-truth regional dishes, ingredients, and flavor profiles; geographic point-of-interest (POI) physical restaurant density metrics extracted via OpenStreetMap (OSM) spatial queries; district-level demographic, household, and urbanization figures from the 2011 Census of India; and empirical consumer ordering preference responses from an academic online food delivery survey. The end-to-end data preparation, regex normalization, schema alignment, median price imputation, and feature assembly pipeline is formally codified and executed within an interactive, reproducible Jupyter Notebook (`data_cleaning_and_transformation.ipynb`).

The algorithmic core of the platform is bifurcated into a Natural Language Processing (NLP) semantic similarity engine and a supervised predictive machine learning regression framework. To evaluate textual menu overlap, the system groups individual dish records into consolidated restaurant documents and computes Term Frequency—

Inverse Document Frequency (TF-IDF) feature vectors across unigram and bigram token spaces (n -gram range 1 to 2). This formulation penalizes ubiquitous culinary stopwords (such as "steamed rice", "plain water", or "roti") while placing strong mathematical weight on distinctive culinary signatures. Intra-locality pairwise Cosine Similarity matrices are subsequently evaluated across all competing kitchens within identical delivery boundaries, yielding a normalized mathematical redundancy coefficient bounded between 0.0 (complete uniqueness) and 1.0 (absolute duplicate). To accommodate user interaction variability in pre-launch testing, a fuzzy token resolution mechanism based on the Levenshtein string distance algorithm is embedded to automatically detect and correct phonetic and typographic misspellings against the canonical culinary vocabulary, while routing genuinely novel dishes into an Out-of-Vocabulary (OOV) novelty scoring pipeline.

For performance forecasting, the system assembles an 11-dimensional feature matrix incorporating competitor density per square kilometer, population scale, intra-locality average menu similarity, planned dish price points (INR), and menu novelty scores. Rigorous 5-Fold Cross-Validation (K_{Fold} , $n = 5$) benchmark experiments were conducted across ensemble regression architectures—including Random Forest Regressors, Extra Trees Regressors, and Gradient Boosting Regressors—evaluated alongside mean and median baseline dummy models under strict target leakage prevention protocols. In predicting monthly delivery order volume, the Gradient Boosting Regressor achieved superior performance with the lowest Mean Absolute Error (MAE = 6,827), outperforming baseline metrics. In forecasting monthly commercial revenue (INR), the Random Forest Regressor emerged as the optimal architecture, yielding an MAE of ₹20.56 Lakhs and an R^2 score of approximately 0.23, effectively capturing complex non-linear interactions between competitive clustering, locality population mass, and pricing sensitivity.

The production platform is deployed as a modular web application powered by Flask, SQLite3, and a high-density Applied Data Science interface. The system features: (1) a dark-themed GIS spatial vector coordinate graph rendering localized kitchen nodes with micro-offsets and dynamic range-slider filtering across similarity thresholds from 50% to 95%; (2) an in-depth trade-area saturation table ranking competitive density, population ratios, and culinary opportunity scores; (3) a dish scarcity versus community demand poll index; and (4) an interactive pre-launch Concept Simulator that evaluates proposed menus, auto-corrects spelling errors, evaluates pricing competitiveness, and outputs an Estimated Performance Score with projected order volume. Verified with a comprehensive automated test suite achieving 10/10 test passes, this system replaces subjective intuition with verifiable mathematical modeling, offering a transformative analytical tool for cloud kitchen entrepreneurs, commercial operators, and food delivery analytics researchers.

2. Introduction:

The contemporary digital economy has fundamentally transformed urban commerce, with the food and beverage industry undergoing one of the most radical structural re-engineerings in modern retail history. Historically, the economic viability of commercial food service establishments was intrinsically tethered to physical location—specifically, high-footfall street retail frontage, visibility along major transportation arteries, visual architectural appeal, and expansive dining room square footage capable of hosting seated patrons. However, the ubiquitous penetration of high-speed mobile telecommunications, app-based on-demand logistics, digital micro-payments, and platform aggregators—most notably Zomato and Swiggy in the Indian domestic ecosystem—has severed the historical link between commercial cooking and brick-and-mortar storefronts. This structural disruption has catalyzed the meteoric rise of **cloud kitchens**, alternatively recognized within industrial literature as ghost kitchens, dark kitchens, virtual kitchens, or delivery-only commissaries.

Cloud kitchens operate as centralized culinary preparation units engineered exclusively for off-premise delivery consumption. By discarding the front-of-house dining infrastructure entirely, kitchen operators eliminate substantial capital liabilities, including premium ground-floor retail leases, complex interior hospitality furnishings, table-service tableware, and extensive front-of-house service personnel such as waiters, hosts, and sommeliers. This lean capital profile theoretically lowers the initial barrier to entry, permitting multi-brand operational models where several disparate culinary concepts operate simultaneously out of a single shared commercial kitchen space. Nevertheless, the elimination of physical storefronts introduces an equally punishing operational reality: commercial survival becomes entirely contingent upon algorithmic visibility and digital shelf competition on aggregator platforms. On an aggregator application, competing kitchens do not compete across city blocks; they compete within a single scrollable digital catalog for consumer attention.

In major metropolitan culinary corridors—such as Koramangala, Indiranagar, and Whitefield in Bangalore, alongside Connaught Place, Saket, and Rajouri Garden in Delhi—dozens of commercial kitchens operate within tightly overlapping 3-to-5 kilometer delivery radiuses. Because delivery platforms prioritize operational proximity to ensure acceptable delivery transit times (typically under 30 to 45 minutes) and preserve food temperature integrity, customer demand within any specific delivery cluster is fundamentally finite and geographically constrained.

In this hyper-dense environment, the overwhelming majority of independent entrepreneurs and emerging multi-brand operators make foundational culinary decisions in an empirical vacuum. Relying predominantly on anecdotal observation, subjective intuition, and crude observations of aggregator "Top Rated" carousels, new

market entrants overwhelmingly default to replicating existing commercial successes. Consequently, delivery corridors witness acute category cannibalization: an influx of near-identical culinary concepts featuring indistinguishable menus dominated by generic Dum Biryani, American-style burgers, commercial pizzas, and standard North Indian butter-gravy curries. Because these competing establishments offer virtually identical item rosters at comparable price points, they aggressively divide a static customer pool, precipitating destructive price-discounting spirals, escalating advertising expenditure on aggregator ad auctions, and driving operational failure rates exceeding sixty percent within the initial six months of commercial launch.

While market corridors are saturated with redundant mass-market cuisines, genuine consumer demand often remains profoundly unmet. Modern metropolitan consumer demographics exhibit highly nuanced dietary and lifestyle requirements—including authentic micro-regional Indian culinary traditions (e.g., coastal Mangalorean, Chettinad, or Kashmiri specialities), health-conscious macro-nutrient profiles (high-protein bowls, keto-friendly options, and vegan formulations), and premium single-portion packaging. Simultaneously, pricing structures are frequently misaligned with localized purchasing power: operators arbitrarily assign menu prices without analyzing the actual median price bands or spending thresholds of households residing within their physical delivery polygon. This structural disconnect between localized supply and localized demand constitutes market whitespace—untapped, commercially viable opportunities where consumer willingness-to-pay exists in the absence of adequate competitive supply.

This project directly addresses this critical information asymmetry by bridging the gap between raw, unstructured culinary catalog data and commercial viability forecasting. Rather than treating restaurant feasibility as a subjective gamble, this study establishes an Applied Data Science (ADS) and Machine Learning framework that transforms digital menu text, spatial infrastructure points, and demographic indicators into actionable, empirical decision metrics.

By utilizing Natural Language Processing (NLP) through unigram and bigram Term Frequency–Inverse Document Frequency (TF-IDF) vectorization, the system mathematically models whole restaurant menus as high-dimensional semantic documents. It executes intra-locality pairwise Cosine Similarity calculations to assign an exact mathematical redundancy coefficient to competing kitchens. Furthermore, the platform fuses this semantic catalog data with real-world physical competitor density derived from OpenStreetMap (OSM) spatial queries and demographic population data from the 2011 Census of India. Finally, through an interactive web workbench powered by Flask, a 40px GIS coordinate vector canvas, and trained regression predictive engines, the system allows prospective operators to simulate proposed menus, correct phonetic typos, evaluate price sensitivity, and predict market performance *before* committing capital to commercial leases, inventory, and kitchen buildouts.

3. Problem Statement:

The transition of food retail toward digital delivery-first business models has unlocked operational agility, yet it introduces profound informational challenges for culinary entrepreneurs. Across major metropolitan hubs, emerging cloud kitchen operators face an alarming commercial failure rate, with a substantial majority forced to close within their initial two quarters of business. These premature business liquidations rarely stem from culinary incompetence or poor product quality. Instead, they result from foundational flaws in pre-launch strategic planning, where operators deploy capital into delivery corridors that cannot sustain additional culinary capacity.

A primary driver of this commercial failure is the absence of computational tools to quantify semantic menu overlap between competing kitchens within identical micro-localities. Currently, entrepreneurs evaluate competition through subjective inspection of consumer-facing food delivery applications. This manual approach fails to detect semantic redundancy across digital catalogs, as competing establishments frequently offer identical preparations under slight naming variations. Because delivery logistics enforce strict geographic radiuses of three to five kilometers to maintain acceptable transit times, competing kitchens with similar menus aggressively divide a static customer base, leading to category cannibalization and destructive price discounting.

This issue is compounded by severe spatial oversaturation and an incomplete understanding of trade-area pricing sensitivity. Commercial aggregator interfaces are engineered for consumer browsing, using recommendation algorithms and customer review scores that obscure the true physical concentration of kitchens in a neighborhood. Operators often assume that high review ratings indicate unmet demand, overlooking that certain urban corridors house excessive preparation units per square kilometer relative to local population. Furthermore, operators assign menu prices based on arbitrary internal markups rather than localized economic benchmarks, failing to align pricing with the median spending thresholds of households within their delivery boundaries.

Ultimately, market entry in the cloud kitchen industry remains dominated by anecdotal hearsay and intuition-driven decision-making. Entrepreneurs observe that staple dishes generate high aggregate platform volume and launch copycat concepts without evaluating whether localized supply has already exceeded consumer absorption capacity. Most critically, the industry lacks an objective, pre-launch simulation environment where an operator can test a proposed menu concept, catalog, and target price point against local competition prior to committing financial capital to commercial leases and equipment. This absence of predictive analytics highlights the urgent need for a data-driven feasibility framework.

4. Existing System:

In conventional commercial practice and academic baseline studies, predictive modeling for restaurant sales and menu comparison relies on lower-hierarchy machine learning architectures, predominantly standalone Decision Trees using greedy recursive binary splitting algorithms like CART, Ordinary Least Squares multiple linear regressions, and raw unweighted bag-of-words or Jaccard similarity indices. The training methodology for these baseline models involves feeding tabular historical transaction records into a single estimator aimed at minimizing residual sum of squares or mean squared error. In such systems, menu text is typically preprocessed through basic whitespace tokenization to construct raw frequency count vectors, while numerical attributes such as average order values, delivery durations, and local competitor counts are passed directly into linear or tree-based regressors without specialized non-linear feature engineering, categorical encoding pipelines, or robust spatial normalization.

This conventional model training framework suffers from severe mathematical, statistical, and architectural limitations that compromise its real-world validity. When training a single Decision Tree, the algorithm evaluates locally optimal feature splits at each interior node using greedy heuristics. Because the tree is trained in isolation without ensemble bagging or gradient-based residual optimization, it exhibits high variance and severe susceptibility to sample noise. In restaurant demand datasets, where trade areas exhibit extreme heterogeneity in population density and purchasing power, a single tree creates rigid, orthogonal decision boundaries that overfit localized training anomalies. In practice, operators rarely tune hyperparameters such as maximum tree depth, minimum samples per leaf, or cost-complexity pruning parameters, leading to over-parameterized trees that essentially memorize the training dataset. A marginal adjustment in an input feature, such as a ten-rupee shift in average dish price or the addition of two nearby competitors, can alter the entire tree topology, leading to volatile and unreliable order volume forecasts. The model sits at the base of the machine learning hierarchy because it lacks the variance-reduction mechanisms of bootstrap aggregated random forests and the sequential error-correction capabilities of gradient boosting algorithms.

Similarly, training multiple linear regression models on food delivery data introduces profound misspecification errors. Linear models operate under the strict assumptions of feature linearity, homoscedasticity, normality of residuals, and independence of predictor variables. However, commercial culinary markets are inherently non-linear and discontinuous. An increase in dish pricing does not produce a uniform linear decline in customer demand; rather, demand remains relatively elastic within a competitive neighborhood price corridor and then collapses abruptly once pricing exceeds local median thresholds. Furthermore, physical competitor density and local population mass exhibit strong multicollinearity, which distorts the estimation of regression coefficients in ordinary least squares models, inflating standard errors and rendering individual feature weights uninterpretable. Additionally, the standard L2 loss

function minimized during ordinary least squares training is disproportionately sensitive to extreme revenue outliers generated by high-volume national chain franchises, pulling regression slopes away from the operating realities of independent, small-scale cloud kitchens.

The training and evaluation of text similarity in existing systems is equally flawed. Relying on raw bag-of-words counts, term co-occurrence matrices, or lexical Jaccard set-similarity fails because these methods operate on binary token presence without inverse document frequency weighting or semantic vector projections. When training on menu corpora, unweighted lexical models treat common items like rice, water, and bread with the same statistical importance as signature culinary preparations, artificially inflating similarity between entirely disparate restaurants while missing genuine semantic duplicates that use varied phrasing. Because raw count vectors exist in high-dimensional, highly sparse spaces without sub-phrase n-gram extraction, Euclidean distance metrics suffer from the curse of dimensionality, where distance measurements between disparate menus collapse toward equidistant values. Furthermore, conventional string matching engines possess no tolerance for typographic misspellings or phonetic variations and completely lack out-of-vocabulary handling, causing analytical pipelines to fail or assign zero relevance whenever a user inputs an innovative culinary fusion or a slightly misspelled regional dish.

These compounding methodological shortcomings ultimately create an insurmountable barrier for entrepreneurs attempting to de-risk their investments before deploying capital. Because current systems rely on unstable single-tree architectures, rigid linear assumptions, unweighted string comparisons, and isolated data silos, they produce feasibility assessments that are both statistically brittle and detached from the physical realities of urban food delivery. The absence of an integrated, pre-launch simulation framework forces operators to treat commercial kitchen launches as high-stakes gambles, entering saturated corridors with duplicate menus and uncalibrated pricing. Finally, conventional model training pipelines rarely implement disciplined k-fold cross-validation or feature-leakage controls. Existing models are frequently trained on static, unstratified train-test splits that inadvertently incorporate post-launch performance signals like customer review counts, driver delivery turnaround times, and promotion-driven platform ratings into the training matrix. Because these operational metrics cannot be known prior to a restaurant opening, their inclusion constitutes severe target leakage, generating artificially inflated cross-validation scores during training that fail entirely during live inference. Coupled with the total exclusion of geospatial point-of-interest infrastructure and census demographic population indicators, existing model training produces blind, uncalibrated estimators that cannot realistically evaluate whether a proposed cloud kitchen concept will survive in a competitive urban trade area.

5. Proposed System:

To overcome the architectural and statistical limitations of conventional baseline models, the proposed Applied Data Science system introduces an end-to-end computational pipeline that operates at the upper tier of the algorithmic hierarchy. In contrast to existing systems that rely on fragile single Decision Trees and rigid linear regressors, the proposed system transitions to advanced ensemble learning architectures, specifically deploying Random Forest Regressors and Gradient Boosting Regressors. By aggregating hundreds of decorrelated decision trees through bootstrap aggregation and sequential residual optimization, these upper-level algorithms dramatically reduce prediction variance, eliminate the overfitting inherent to single-tree architectures, and successfully model complex, non-linear interactions between competitive density, population mass, and pricing thresholds. The training pipeline enforces strict 5-Fold Cross-Validation across five stratified subsets alongside automated column transformer preprocessing, which applies one-hot encoding to categorical features, median imputation to missing numerical values, and standard scaling across continuous variables. Furthermore, the model eliminates target leakage by training strictly on pre-launch parameters, ensuring that the resulting estimators provide realistic, generalizable performance forecasts for unlaunched culinary concepts.

At the natural language processing layer, the proposed framework replaces rudimentary bag-of-words counts and keyword filtering with an n-gram Term Frequency–Inverse Document Frequency vectorization pipeline spanning unigram and bigram token spaces. This mathematical formulation computes localized term frequencies while downweighting ubiquitous culinary terms through inverse document frequency scaling, ensuring that signature dishes receive high feature salience while common staples do not distort similarity scores. Rather than computing brittle set intersections, the system projects restaurant menus into high-dimensional vector spaces and evaluates intra-locality pairwise Cosine Similarity matrices across all competing kitchens within identical delivery radiuses. Because the cosine metric divides the inner product of menu vectors by their Euclidean norms, it provides length-invariant normalization, allowing small ten-item menus and extensive sixty-item catalogs to be compared without document-length bias. To handle noisy user inputs during live inference, the system incorporates an approximate string matching layer based on the Levenshtein distance algorithm, which evaluates character edit distances against a canonical taxonomy derived from the Indian Food 101 database to automatically correct typographic errors, while simultaneously routing unrecognized fusion items into an out-of-vocabulary novelty scoring pipeline.

This natural language intelligence is fully integrated with a multi-source geospatial and demographic feature store that grounds abstract text similarities in physical market realities. The system merges digital menu listings with spatial point-of-interest competitor counts extracted from OpenStreetMap, district-level demographic and household indicators from the 2011 Census of India, and consumer dining preference

metrics from academic delivery surveys. By assigning deterministic micro-coordinate offsets to individual kitchen nodes, the platform renders a technical coordinate vector space where users can interactively visualize trade-area clustering and dynamically filter menu redundancy vectors across customizable similarity thresholds ranging from fifty to ninety-five percent. The culmination of this architecture is an interactive pre-launch Concept Simulator that ingests a candidate dish list, planned price point, and target neighborhood to evaluate localized competition, identify unmet culinary whitespace, and generate an Estimated Performance Prediction with projected monthly order volumes and revenue bounds. Entirely codified and reproducible through an automated six-stage Jupyter Notebook pipeline, the proposed system elevates feasibility analysis from subjective intuition to an empirical, mathematically grounded decision-support engine.

At the predictive modeling tier, the proposed system advances beyond baseline single-estimator models by deploying ensemble architectures, specifically Random Forest and Gradient Boosting regressors. By aggregating hundreds of decorrelated decision trees and sequentially correcting residual errors, these upper-level algorithms drastically reduce model variance and eliminate the severe overfitting characteristic of single decision trees. Validated through disciplined 5-Fold Cross-Validation on leakage-free pre-launch parameters, the models accurately capture complex, non-linear relationships between trade-area competitor density, demographic population, and menu pricing to output robust forecasts for monthly orders and revenue.

To uncover untapped commercial opportunities, the system introduces a dedicated whitespace identification algorithm that balances culinary catalog scarcity against empirical consumer demand signals. By cross-referencing local consumer survey preferences with active competitor offerings, the platform calculates an objective opportunity score for each dish category within a delivery polygon. This mechanism works alongside an n-gram TF-IDF vectorizer and a Levenshtein-based fuzzy matching module, ensuring that menu items are normalized against a canonical culinary taxonomy while identifying genuine dietary gaps that existing delivery kitchens have failed to serve.

The complete analytical architecture is operationalized as an interactive web workbench built using Python, Flask, and an optimized SQLite database. The client interface features a technical GIS coordinate map where users can visualize micro-trade areas and dynamically filter menu redundancy links using a variable similarity threshold slider ranging from fifty to ninety-five percent. Coupled with detailed neighborhood saturation tables and a real-time Concept Simulator, the platform empowers culinary entrepreneurs to test proposed menus and price points in an empirical sandbox prior to committing capital to commercial facility leases.

6. About Dataset:

The system integrates five distinct empirical datasets into a unified analytical feature store:

Dataset Name	Source & Format	Records & Features	Primary Role in System
Zomato Menus (cleaned_zomato_menus.csv)	Web-scraped catalog (CSV)	4,000+ dish listings across Bangalore & Delhi	Baseline dish catalog, restaurant names, price bands, ratings, and locations.
Indian Food 101 (indian_food_101.csv)	Culinary Database (CSV)	255 curated dishes with ingredients & diet	Ground truth dictionary for fuzzy spelling correction & canonical dish mapping.
OSM Density (osm_restaurant_density.csv)	OpenStreetMap API (CSV)	Locality-level POI spatial density counts	Physical commercial competitor density per square kilometer.
Census 2011 (census_districts_2011.csv)	Census of India (CSV)	District demographics & household counts	Trade-area population scale & per-capita kitchen penetration index.
Online Delivery Survey (online_food_delivery_survey.csv)	Survey dataset (CSV)	Consumer dining preferences & price tiers	Weights consumer demand signals against locality supply gaps.

Jupyter Notebook Transformation Stages (data_cleaning_and_transformation.ipynb):

- **Stage 1 (Ingestion):** Automated schema verification, column casting, null-value audits.
- **Stage 2 (Cleaning & Aggregation):** Regex punctuation cleaning; median price imputation; grouping individual dish records by [Restaurant, City, Location] into consolidated restaurant document strings.
- **Stage 3 (Normalization & Joins):** Trade locality standardization; OSM density join; Census population join; deterministic micro-coordinate offset generation for GIS spatial rendering.
- **Stage 4 (TF-IDF & Cosine Similarity):** Fitting TfidfVectorizer(ngram_range=(1,2)); computing pairwise similarity matrices within each locality cluster.
- **Stage 5 (Feature Assembly & Split):** Compiling model_features.csv (11 engineering features); 80/20 train/test split; fitting StandardScaler.
- **Stage 6 (Artifact Persistence):** Writing sanitized SQLite tables (menus.sqlite3), CSV files, and logging execution metrics to join_audit.json.

7. Technologies Used:


Category	Technologies / Libraries	Purpose in System Architecture
Programming Language	Python 3.x	Core runtime environment for backend, analytical models, and notebooks.
Machine Learning & NLP	Scikit-learn, NumPy, Pandas, FuzzyWuzzy	TF-IDF vectorization, pairwise Cosine Similarity, Ridge regression, StandardScaler, Levenshtein fuzzy string distance.
Notebook Environment	Jupyter Notebook (nbclient, nbformat)	Reproducible execution and interactive documentation of data cleaning and transformation.
Backend Web Framework	Flask, SQLite3, Werkzeug	REST API routing, user session authentication, and transactional database storage.
Frontend & UI	HTML5, CSS3 (Slate Clean Tech), JavaScript (ES6+)	Clean Applied Data Science interface: GIS coordinate grid, dynamic slider controls, and responsive telemetry.

8. Data Flow Diagram:

The system follows a modular 4-tier data and prediction pipeline:



Level 0: Raw Data Sources (Context Diagram)



Level 1: Feature Extraction & NLP Processing

E1 Zomato Menus Catalog	E2 Indian Food 101 Taxonomy (255 dishes)
----------------------------	--

Menu listings

Canonical taxonomy



1.1
DATA CLEANING & IMPUTATION Regex text cleaning, lowercasing, fuzzy dish-name correction, median price imputation

Cleaned dish records



1.2
RESTAURANT DOCUMENT AGGREGATION Groups dishes by [Restaurant, City, Locality] into one text document per restaurant

Restaurant documents



1.3
n-GRAM (1-2) TF-IDF VECTORIZER Unigram + bigram TF-IDF weighting of each restaurant document

Sparse TF-IDF matrix



1.4
INTRA-LOCALITY PAIRWISE COSINE SIMILARITY ENGINE Pairwise cosine similarity matrix within each locality, scores 0.0 to 1.0

Menu redundancy scores (0.0-1.0)



E3 OSM Competitor Density	E4 Census of India 2011	E5 Online Delivery Survey
---------------------------------	-------------------------------	---------------------------------

Competitor density

Population mass

Price sensitivity




1.5
SPATIAL POI & CENSUS DEMOGRAPHIC NORMALIZATION Scales competitor density per sq. km, population mass and survey price-sensitivity signals

Normalized spatial features



2.1
11-FEATURE MATRIX ASSEMBLY (continues in Level 2) Receives redundancy scores and normalized spatial features

Level 2: Model Training & Storage Layer



Level 3: Application & Analytics Workbench

E6

Cloud Kitchen Operator / Market Analyst (Browser Client)

▲
Slider events / map view
▼

3.2a

INTERACTIVE GIS
COORDINATE VECTOR MAP
40px grid, 50%-95% dynamic overlap
slider

▲
GET /api/insights, /api/summary (JSON)
▼

▲
Locality selection / tables
▼

3.2b

TRADE-AREA SATURATION &
DISH WHITESPACE TABLES
Koramangala, Indiranagar, etc.

▲
GET /api/analysis/areas (JSON)
▼

▲
Concept input / viability score
▼

3.2c

PRE-LAUNCH CONCEPT
SIMULATOR
Levenshtein typo correction, 0-100
viability score

▲
POST /api/predict-concept (request / JSON)
▼

3.1

FLASK REST API GATEWAY

GET /api/summary | GET /api/insights | GET /api/analysis/areas | POST /api/predict-concept
Routes requests, queries the database, loads the pickled models and returns JSON

▲
Features, scores
|

D2

menus.sqlite3

▲
Fitted models
|

D3

Model pickles (.pkl)



▲
Audit log & status
|

D4

join_audit.json

1. Output:

The screenshot displays the Bitewise local menu dashboard. At the top, there's a dark header with the Bitewise logo and a 'Sign In' button. The main content area is light blue and features a central white sign-in form. The form has a 'WELCOME TO BITEWISE' message, a 'Sign in to Bitewise' title, and instructions to enter credentials for neighborhood competition and menu ideas. It includes 'Sign In' and 'Register' links, a 'USERNAME' field with the example 'e.g. analyst_user', a 'PASSWORD' field with a 'Minimum 6 characters' requirement, and a 'Sign In to Bitewise' button. Below the form, there's a link for 'Need access? Create new account'. At the bottom, a 'DEMO ACCOUNTS' section lists two roles: 'KITCHEN OWNER' and 'DATA ANALYST'.



MARKET & MENU ANALYSIS

Active Delivery Areas

Neighborhood Competition & Food Opportunities

Analyze delivery neighborhoods across Bangalore and Delhi to see how crowded they are, what dishes are most common, and where customers want more food choices.

Test in Menu Simulator

Back to Map Overview

DELIVERY AREAS

25

Markets included in the data

TRACKED KITCHENS

200

Restaurants included in the data

UNIQUE DISHES

137

Cataloged menu offerings

LESS-COMMON DISHES

12

Dishes found on fewer local menus

Select City:

All Cities

All Categories

Main Course

Appetizers

Snacks

Desserts

Beverages

Customer Polls

Search dish (e.g. Biryani)...

AREA COMPARISON

Menu similarity and restaurant count

25 delivery areas

A higher percentage means menus in that area are more alike. The count shows restaurants and other food places found in the available data.

Indiranagar

BANGALORE

185 tracked kitchens • 50 local food places • ₹293 avg dish

Gobi 65 (8) • Shenanigans Chicken (8) • Vegetable Biryani (8)

Diverse Menu

14%

Jayanagar

BANGALORE

187 tracked kitchens • 40 local food places • ₹290 avg dish

Chicken 65 (12) • Shahi Paneer (8) • Chilli Paneer Dry (8)

Diverse Menu

14%

Koramangala

BANGALORE

175 tracked kitchens • 26 local food places • ₹309 avg dish

Chole Bhature (10) • Fish Tikka (10) • Jalebi with Rabri (10)

Diverse Menu

15%

Mg Road

BANGALORE

179 tracked kitchens • 34 local food places • ₹289 avg dish

Chilli Paneer Dry (8) • Jalebi with Rabri (8) • Curd Rice (8)

Diverse Menu

13%

DISH COVERAGE

Common Dishes & Food Opportunities

137 dishes

Dishes found on fewer menus may be worth investigating. Click any card to test it in the menu checker.

Korean Fried Chicken

100% Wanted (2 votes)

Market Whitespace

Community Voted / High Demand • ₹280 avg • 0% presence (0/200 kitchens)

Less common dish

Test in Simulator

Truffle Mushroom Pizza

100% Wanted (2 votes)

Market Whitespace

Community Voted / High Demand • ₹280 avg • 0% presence (0/200 kitchens)

Less common dish


Test in Simulator

Burger

67% Wanted (2 votes)

Market Whitespace

Community Voted / High Demand • ₹280 avg • 0% presence (0/200 kitchens)



Bitewise

Bitewise local menu dashboard

Overview

Neighborhoods

Menu Simulator

How It Works

Log out

TEST A MENU IDEA

Menu checker ready

Check your menu before launch

Choose an area, enter the dishes you plan to sell, and see how often they already appear nearby. You can also compare your price with the local average.

Explore Neighborhoods

Back to Overview

ENTER YOUR MENU

Your proposed dishes

Choose city

Bangalore

Neighbourhood / delivery area

Indiranagar

Dishes you plan to sell (one dish per line)

Chicken Biryani
Falafel Wrap
Truffle Fries

Price per dish (₹) (optional)

250

Check this menu →

MENU CHECK RESULTS

Enter your dishes to begin

Enter your dishes above and click "Check this menu" to see local menu similarity, coverage, and an explainable research score.

LOCAL CUSTOMER FEEDBACK

What would you try?

Share a quick response for a dish or see feedback already collected for this area.

Enter dish name (e.g. Korean Fried Chicken, Avocado Toast)


I would try it

Not for me

Customer Poll Demand in Indiranagar, Bangalore

Click any dish to add to your menu above

Truffle Mushroom Pizza



2. Conclusion:

The Cloud Kitchen Menu Redundancy & Trade-Area Whitespace Prediction System successfully demonstrates that modern Applied Data Science and Machine Learning methodologies can systematically de-risk digital food delivery entrepreneurship. By substituting speculative intuition and anecdotal heuristics with mathematical vector-space similarity matrices, multi-source spatial demographic integration, and ensemble predictive modeling, the platform establishes an empirical foundation for commercial culinary feasibility analysis.

The primary technical contribution of this research lies in its multi-layered analytical pipeline. At the linguistic tier, formulating restaurant menus through n -gram Term Frequency–Inverse Document Frequency (TF-IDF) vectorization and length-normalized pairwise Cosine Similarity successfully resolves the challenge of unquantified catalog overlap, allowing the system to identify duplicate culinary concepts even when items are marketed under disparate brand names. Concurrently, embedding the Levenshtein string distance algorithm provides robust typo tolerance against ground-truth culinary taxonomies while establishing an automated routing mechanism for novel culinary innovations. At the macroeconomic and spatial tier, fusing digital menu catalogs with OpenStreetMap competitor densities and Census demographic indicators ensures that semantic similarities are grounded in the physical realities of urban delivery corridors. Furthermore, empirical benchmarking demonstrated that upper-hierarchy ensemble architectures—specifically Gradient Boosting Regressors for order volume prediction and Random Forest Regressors for revenue forecasting—consistently outperform baseline linear models and single

decision trees by capturing complex non-linear interactions across competitive density, population scale, and localized pricing thresholds.

From a practical implementation standpoint, the deployment of an interactive, clean Applied Data Science workbench—complete with a dynamic GIS coordinate vector canvas, trade-area saturation metrics, and a pre-launch Concept Simulator—provides prospective operators with an actionable testing sandbox. Entrepreneurs can simulate candidate menus, optimize price points relative to neighborhood median purchasing power, and identify high-demand whitespace opportunities prior to leasing commercial preparation units and deploying financial capital. The entire data engineering and transformation workflow remains fully documented and reproducible through an automated six-stage Jupyter Notebook pipeline, establishing a robust computational blueprint for future academic research in hyper-local commercial intelligence.

- **Limitations and Future Work:**

While the system achieves high predictive stability, several prospective avenues remain for ongoing research and engineering refinement:

- **Dynamic Sentiment and Review Mining:**

The current architecture evaluates static menu catalogs and survey benchmarks. Future iterations will integrate real-time web-scraping pipelines to ingest live customer reviews and social sentiment streams, employing aspect-based sentiment analysis (ABSA) via transformer models (such as BERT) to gauge qualitative consumer dissatisfaction with incumbent offerings.

Road-Network Isochrone

- **Modeling:**

Physical delivery boundaries are currently approximated using localized geographic coordinates. Enhancing the spatial engine with open routing graph services (such as OSRM) will allow the system to compute precise travel-time isochrones based on real-time traffic conditions and road network topologies rather than Euclidean buffers.

- **Supply-Chain Cost Indexing:**

Integrating dynamic wholesale commodity and ingredient pricing APIs will enable the concept simulator to perform real-time gross margin and contribution margin forecasting alongside demand predictions.

- **Multi-Metropolitan Scaling:**

Expanding the pre-trained feature stores beyond Bangalore and Delhi into Tier-2 urban consumption corridors will broaden the demographic variance of the regression models and validate cross-regional culinary transferability.